

# A Good Listner

---

Kim Takala, Mai Nguyen, Matteo Murdaca

August 12, 2025

Fazer - the better Milky Way

**Online meetings have a problem**

---

# The problem

- Too much time is spent transcribing and summarizing meetings.
- The secretary is unable to actively participate in the meeting because of the need to write down minutes.

## Solution

---

*A Good Listener* can transcribe meetings recordings, differentiate speakers, extract timestamps, and summarize highlights. It can also convert document pictures into text, translate transcriptions and summary, and answer questions about the meeting.

# Core Features

- Device agnostic
- Audio retrieval
- Audio transcribing
- Optical Character Recognition to digitize documents
- AI assistant for tasks such as summarizing, translating, and information retrieval

- **Data:** Audio files from starter code and documents (e.g., PDF, PPTX, PNG)
- **Tools:** Docker, WhisperX, Gemma3, MiniCMP, Ollama

# Users

---



## Target Users

- Office workers who attend many online meetings.
- Reduces time and effort in keeping track of meeting content.

# User Interaction

1. Run `'podman run assistant'`.
2. The assistant asks for the recordings directory.
3. Automatically summarizes the meeting.
4. Offers translation of the summary.
5. Allows further questions to retrieve meeting information.

It's finally demo time!

# Impact

---

# Impact

- Fully local deployment to allow for maximum privacy.
- Reduces extra effort to write down meeting checkpoints, enhancing productivity.

**How to improve?**

---

## Future Improvements

- Assigning speaker names to speaker IDs for better interaction.
- Saving previous sessions as updated context for the bot.
- Implementing a graphical interface to improve usability.

**Thank You!**

Thank you for your attention!